

Multiple Linear Regression

Importing libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

Importing dataset

```
In [2]: dataset = pd.read_csv('50_Startups-2.csv')
X = dataset.iloc[ : , : -1].values
y = dataset.iloc[ : , -1].values
```

```
In [3]: print(X)
print("\n\n")
print(y)
```

```
[[165349.2 136897.8 471784.1 'New York']  
[162597.7 151377.59 443898.53 'California']  
[153441.51 101145.55 407934.54 'Florida']  
[144372.41 118671.85 383199.62 'New York']  
[142107.34 91391.77 366168.42 'Florida']  
[131876.9 99814.71 362861.36 'New York']  
[134615.46 147198.87 127716.82 'California']  
[130298.13 145530.06 323876.68 'Florida']  
[120542.52 148718.95 311613.29 'New York']  
[123334.88 108679.17 304981.62 'California']  
[101913.08 110594.11 229160.95 'Florida']  
[100671.96 91790.61 249744.55 'California']  
[93863.75 127320.38 249839.44 'Florida']  
[91992.39 135495.07 252664.93 'California']  
[119943.24 156547.42 256512.92 'Florida']  
[114523.61 122616.84 261776.23 'New York']  
[78013.11 121597.55 264346.06 'California']  
[94657.16 145077.58 282574.31 'New York']  
[91749.16 114175.79 294919.57 'Florida']  
[86419.7 153514.11 0.0 'New York']  
[76253.86 113867.3 298664.47 'California']  
[78389.47 153773.43 299737.29 'New York']  
[73994.56 122782.75 303319.26 'Florida']  
[67532.53 105751.03 304768.73 'Florida']  
[77044.01 99281.34 140574.81 'New York']  
[64664.71 139553.16 137962.62 'California']  
[75328.87 144135.98 134050.07 'Florida']  
[72107.6 127864.55 353183.81 'New York']  
[66051.52 182645.56 118148.2 'Florida']  
[65605.48 153032.06 107138.38 'New York']  
[61994.48 115641.28 91131.24 'Florida']  
[61136.38 152701.92 88218.23 'New York']  
[63408.86 129219.61 46085.25 'California']  
[55493.95 103057.49 214634.81 'Florida']  
[46426.07 157693.92 210797.67 'California']  
[46014.02 85047.44 205517.64 'New York']  
[28663.76 127056.21 201126.82 'Florida']  
[44069.95 51283.14 197029.42 'California']  
[20229.59 65947.93 185265.1 'New York']  
[38558.51 82982.09 174999.3 'California']  
[28754.33 118546.05 172795.67 'California']  
[27892.92 84710.77 164470.71 'Florida']  
[23640.93 96189.63 148001.11 'California']  
[15505.73 127382.3 35534.17 'New York']  
[22177.74 154806.14 28334.72 'California']  
[1000.23 124153.04 1903.93 'New York']  
[1315.46 115816.21 297114.46 'Florida']  
[0.0 135426.92 0.0 'California']  
[542.05 51743.15 0.0 'New York']  
[0.0 116983.8 45173.06 'California']]
```

```
[192261.83 191792.06 191050.39 182901.99 166187.94 156991.12 156122.51  
155752.6 152211.77 149759.96 146121.95 144259.4 141585.52 134307.35  
132602.65 129917.04 126992.93 125370.37 124266.9 122776.86 118474.03  
111313.02 110352.25 108733.99 108552.04 107404.34 105733.54 105008.31  
103282.38 101004.64 99937.59 97483.56 97427.84 96778.92 96712.8  
96479.51 90708.19 89949.14 81229.06 81005.76 78239.91 77798.83
```

```
71498.49 69758.98 65200.33 64926.08 49490.75 42559.73 35673.41  
14681.4 ]
```

Encoding categorical data

```
In [4]: from sklearn.compose import ColumnTransformer  
from sklearn.preprocessing import OneHotEncoder  
  
column_transformer = ColumnTransformer(transformers = [('encoder',  
                                                    OneHotEncoder(), [3]]),  
                                       remainder = "passthrough")  
X = np.array(column_transformer.fit_transform(X))
```

```
In [5]: print("Encoded X is:")  
print("-----")  
print(X)
```

Encoded X is:

```
-----
[[0.0 0.0 1.0 165349.2 136897.8 471784.1]
 [1.0 0.0 0.0 162597.7 151377.59 443898.53]
 [0.0 1.0 0.0 153441.51 101145.55 407934.54]
 [0.0 0.0 1.0 144372.41 118671.85 383199.62]
 [0.0 1.0 0.0 142107.34 91391.77 366168.42]
 [0.0 0.0 1.0 131876.9 99814.71 362861.36]
 [1.0 0.0 0.0 134615.46 147198.87 127716.82]
 [0.0 1.0 0.0 130298.13 145530.06 323876.68]
 [0.0 0.0 1.0 120542.52 148718.95 311613.29]
 [1.0 0.0 0.0 123334.88 108679.17 304981.62]
 [0.0 1.0 0.0 101913.08 110594.11 229160.95]
 [1.0 0.0 0.0 100671.96 91790.61 249744.55]
 [0.0 1.0 0.0 93863.75 127320.38 249839.44]
 [1.0 0.0 0.0 91992.39 135495.07 252664.93]
 [0.0 1.0 0.0 119943.24 156547.42 256512.92]
 [0.0 0.0 1.0 114523.61 122616.84 261776.23]
 [1.0 0.0 0.0 78013.11 121597.55 264346.06]
 [0.0 0.0 1.0 94657.16 145077.58 282574.31]
 [0.0 1.0 0.0 91749.16 114175.79 294919.57]
 [0.0 0.0 1.0 86419.7 153514.11 0.0]
 [1.0 0.0 0.0 76253.86 113867.3 298664.47]
 [0.0 0.0 1.0 78389.47 153773.43 299737.29]
 [0.0 1.0 0.0 73994.56 122782.75 303319.26]
 [0.0 1.0 0.0 67532.53 105751.03 304768.73]
 [0.0 0.0 1.0 77044.01 99281.34 140574.81]
 [1.0 0.0 0.0 64664.71 139553.16 137962.62]
 [0.0 1.0 0.0 75328.87 144135.98 134050.07]
 [0.0 0.0 1.0 72107.6 127864.55 353183.81]
 [0.0 1.0 0.0 66051.52 182645.56 118148.2]
 [0.0 0.0 1.0 65605.48 153032.06 107138.38]
 [0.0 1.0 0.0 61994.48 115641.28 91131.24]
 [0.0 0.0 1.0 61136.38 152701.92 88218.23]
 [1.0 0.0 0.0 63408.86 129219.61 46085.25]
 [0.0 1.0 0.0 55493.95 103057.49 214634.81]
 [1.0 0.0 0.0 46426.07 157693.92 210797.67]
 [0.0 0.0 1.0 46014.02 85047.44 205517.64]
 [0.0 1.0 0.0 28663.76 127056.21 201126.82]
 [1.0 0.0 0.0 44069.95 51283.14 197029.42]
 [0.0 0.0 1.0 20229.59 65947.93 185265.1]
 [1.0 0.0 0.0 38558.51 82982.09 174999.3]
 [1.0 0.0 0.0 28754.33 118546.05 172795.67]
 [0.0 1.0 0.0 27892.92 84710.77 164470.71]
 [1.0 0.0 0.0 23640.93 96189.63 148001.11]
 [0.0 0.0 1.0 15505.73 127382.3 35534.17]
 [1.0 0.0 0.0 22177.74 154806.14 28334.72]
 [0.0 0.0 1.0 1000.23 124153.04 1903.93]
 [0.0 1.0 0.0 1315.46 115816.21 297114.46]
 [1.0 0.0 0.0 0.0 135426.92 0.0]
 [0.0 0.0 1.0 542.05 51743.15 0.0]
 [1.0 0.0 0.0 0.0 116983.8 45173.06]]
```

Splitting the dataset into Training and testing dataset

```
In [6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
                                                    random_state = 0)
```

Training the Model

```
In [7]: from sklearn.linear_model import LinearRegression
linear_regression = LinearRegression()
linear_regression.fit(X_train, y_train)
```

```
Out[7]: LinearRegression()
```

Predicting the values

```
In [8]: y_pred = linear_regression.predict(X_test)
print("Predicted values:")
print("y_test\tY_pred")

np.set_printoptions(precision=2)
print(np.concatenate((y_test.reshape(len(y_test), 1),
                      y_pred.reshape(len(y_pred), 1)),
                      1))
```

Predicted values:

y_test	Y_pred
[103282.38	103015.2]
[144259.4	132582.28]
[146121.95	132447.74]
[77798.83	71976.1]
[191050.39	178537.48]
[105008.31	116161.24]
[81229.06	67851.69]
[97483.56	98791.73]
[110352.25	113969.44]
[166187.94	167921.07]]